

Sensorless PMSM Drive — Theory, Observer Design, and Parameters

This document is the standalone technical reference for the simulation built across `01_pmsm_machine_and_controller.ipynb`, `02_pmsm_with_parallel_observer.ipynb`, and `03_pmsm_sensorless_control.ipynb`. It collects the machine model, the current-control design, the sensorless position/speed observer, the exact gain formulas used, and the full parameter set — with the reasoning behind each choice — in one place.

Contents

- 1. [System overview](#)
 - 2. [Machine model \(dq frame\)](#)
 - 3. [Mechanical model and load](#)
 - 4. [Current control](#)
 - 5. [Sensorless position/speed estimation](#)
 - 6. [Model and gain parameters \(reference table\)](#)
 - 7. [Known limitations and observed behavior](#)
 - 8. [References](#)
-

1. System overview

The plant is a permanent-magnet synchronous motor (PMSM) driving a fan, which loads it with a torque that grows with the square of speed. The motor is torque-controlled: an external torque reference `Tref(t)` steps through the schedule `[0, 10, 30, 50, 70, 10, 70, -20, 50]` N·m, holding each level until the drive settles before moving to the next.

The simulation is built in three stages, each adding exactly one piece so that the specific cost of each addition is isolated:

Stage	Notebook	Controller feedback	Observer
1	<code>01_...ipynb</code>	plant's own real θ_r , ω_{mech}	none
2	<code>02_...ipynb</code>	still real θ_r , ω_{mech}	runs in parallel, passively monitored
3	<code>03_...ipynb</code>	observer's estimated $\hat{\theta}$, $\hat{\omega}$	closes the loop (fully sensorless)

Three physical/mathematical pieces make up the whole system:

- a **dq-frame machine model** (Section 2) — the "real" plant, always evaluated with its own true rotor angle θ_r , regardless of what the controller believes;
- **dq PI current controllers with EMF feedforward** (Section 4) — the torque-mode control loop, tuned against an assumed inverter response delay;
- an **$\alpha\beta$ -frame back-EMF observer + PLL** (Section 5) — estimates $\hat{\theta}$, $\hat{\omega}$ from measurable stator currents and applied voltage alone, with no mechanical sensor.

2. Machine model (dq frame)

2.1 Convention: power-invariant Park transform, motor sign convention

The machine-model structure is adapted from the source reference document "RL_Modelo PMSM.pdf", which derives a PMSM model for a **generator**, using the **power-invariant** Blondel-Park transform:

$$T_{BP}(\theta) = \sqrt{2/3} \cdot \begin{bmatrix} \cos\theta & -\sin\theta & 1/\sqrt{2} \\ \cos(\theta-2\pi/3) & -\sin(\theta-2\pi/3) & 1/\sqrt{2} \\ \cos(\theta+2\pi/3) & -\sin(\theta+2\pi/3) & 1/\sqrt{2} \end{bmatrix}$$

(orthogonal, so $T_{BP}^{-1} = T_{BP}^T$). Two adjustments were made to use this source for our motor:

(a) Sign convention. The source's stator KVL is written for current flowing *out* of the machine (generator reference): $u = -R \cdot i - L \cdot di/dt + e$. Substituting $i_{\text{motor}} = -i_{\text{gen}}$ converts this to the standard motor KVL, current flowing *in*: $u = R \cdot i + L \cdot di/dt + e$.

(b) Where the power-invariant factor shows up. Because the transform is power-invariant rather than the more common amplitude-invariant (conventional $2/3$) transform, a factor of $\sqrt{3/2}$ appears explicitly wherever the permanent-magnet flux ψ_m couples into an equation — voltage, torque, and the $\alpha\beta$ back-EMF used by the observer — while R_s , L_d , L_q and the $dq \leftrightarrow \alpha\beta$ relationship (a plain rotation, since discrete phases are never modeled here) are unaffected. This was **verified by power balance** (not just carried over from the source), see §2.3.

2.2 Electrical and torque equations

With $\omega_r = p \cdot \omega_{\text{mech}}$ the electrical angular speed and θ_r the electrical rotor angle:

$$\begin{aligned} u_{sd} &= R_s \cdot i_{sd} + L_d \cdot di_{sd}/dt - \omega_r \cdot L_q \cdot i_{sq} \\ u_{sq} &= R_s \cdot i_{sq} + L_q \cdot di_{sq}/dt + \omega_r \cdot L_d \cdot i_{sd} + \sqrt{3/2} \cdot \omega_r \cdot \psi_m \\ T_e &= p \cdot \sqrt{3/2} \cdot \psi_m \cdot i_{sq} + p \cdot (L_d - L_q) \cdot i_{sd} \cdot i_{sq} \end{aligned}$$

Define $\psi_{m,pi} = \sqrt{3/2} \cdot \psi_m$ (the power-invariant flux constant, §6) so these read more simply:

$$\begin{aligned} u_{sq} &= R_s \cdot i_{sq} + L_q \cdot di_{sq}/dt + \omega_r \cdot L_d \cdot i_{sd} + \omega_r \cdot \psi_{m,pi} \\ T_e &= p \cdot \psi_{m,pi} \cdot i_{sq} + p \cdot (L_d - L_q) \cdot i_{sd} \cdot i_{sq} \end{aligned}$$

2.3 Power-balance verification

Multiplying the voltage equations by their respective currents and summing:

$$\begin{aligned} P &= u_{sd} \cdot i_{sd} + u_{sq} \cdot i_{sq} \\ &= R_s \cdot (i_{sd}^2 + i_{sq}^2) && \text{(copper loss)} \\ &+ \frac{d}{dt} \left[\frac{1}{2} (L_d \cdot i_{sd}^2 + L_q \cdot i_{sq}^2) \right] && \text{(stored magnetic energy rate)} \\ &+ \omega_r \cdot [(L_d - L_q) \cdot i_{sd} \cdot i_{sq} + \psi_{m,pi} \cdot i_{sq}] && \text{(electromagnetic power)} \end{aligned}$$

The bracketed electromagnetic-power term, divided by p (i.e. multiplied by $\omega_r/p = \omega_{mech}$), equals $T_e \cdot \omega_{mech}$ exactly — confirming the equations (and the placement of $\sqrt{3/2}$) are self-consistent: $P = I^2 R_{loss} + \text{rate of stored energy} + \text{mechanical power out}$. Because the transform is power-invariant, P needs no extra $3/2$ scaling factor (unlike the conventional transform, where $P = (3/2)(u_{sd} \cdot i_{sd} + u_{sq} \cdot i_{sq})$).

2.4 Reference-frame relations (used by the observer/controller, Section 5)

Since discrete phases (abc) are never modeled, the only relation needed between the dq frame and the stationary $\alpha\beta$ frame is a plain rotation by the (real or estimated) angle:

$$\begin{aligned} i_{\alpha} &= i_d \cdot \cos\theta - i_q \cdot \sin\theta & i_{\beta} &= i_d \cdot \sin\theta + i_q \cdot \cos\theta & (dq \rightarrow \alpha\beta) \\ i_d &= i_{\alpha} \cdot \cos\theta + i_{\beta} \cdot \sin\theta & i_q &= -i_{\alpha} \cdot \sin\theta + i_{\beta} \cdot \cos\theta & (\alpha\beta \rightarrow dq) \end{aligned}$$

(and identically for voltage). Ground-truth $\alpha\beta$ back-EMF, used for validation:

$$\begin{aligned} e_{\alpha} &= \sqrt{3/2} \cdot \omega_r \cdot \psi_{m,pi} \cdot \cos\theta_r = \omega_r \cdot \psi_{m,pi} \cdot \cos\theta_r \\ e_{\beta} &= \omega_r \cdot \psi_{m,pi} \cdot \sin\theta_r \end{aligned}$$

3. Mechanical model and load

$$J \cdot d\omega_{mech}/dt = T_e - T_{load}(\omega_{mech}) - B \cdot \omega_{mech} \qquad d\theta_r/dt = \omega_r = p \cdot \omega_{mech}$$

$B = 0$ (no viscous friction assumed — only the fan load opposes rotation). The fan load is quadratic in speed and opposes rotation in *either* direction, continuous through zero:

$$T_{load}(\omega_{mech}) = k_{fan} \cdot \omega_{mech} \cdot |\omega_{mech}|$$

k_{fan} is set from the one operating point given: 11 kW at 1500 rpm. With $\omega_{mech,rated} = 1500 \cdot 2\pi/60 \approx 157.08$ rad/s and $T_{rated} = P_{rated}/\omega_{mech,rated} \approx 70.03$ N·m (matching the largest value in the torque schedule — confirming the schedule's units are N·m):

$$k_{fan} = T_{rated} / \omega_{mech,rated}^2 \approx 2.838 \times 10^{-3} \text{ N·m/(rad/s)}^2$$

$J = 0.05 \text{ kg}\cdot\text{m}^2$ is a user-selected value (not given in the literature) representative of an 11 kW industrial PMSM rotor + small fan load.

A note on the $-20 \text{ N}\cdot\text{m}$ schedule step: since T_{load} opposes rotation in either direction, this doesn't fail to be "balanced" — the drive simply reverses and settles at a negative-speed equilibrium $\omega_{\text{mech}} \cdot |\omega_{\text{mech}}| = -20/k_{\text{fan}} \Rightarrow \omega_{\text{mech}} \approx -83.9 \text{ rad/s}$, where both T_e and ω_{mech} are negative, so mechanical power $T_e \cdot \omega_{\text{mech}}$ stays positive throughout (motoring, not braking). This step was kept deliberately as a stress case — see Section 7.

4. Current control

4.1 References ($i_d^* = 0$ strategy)

L_d and L_q differ by only $\sim 2.2\%$ (see §6), so the reluctance-torque term $p(L_d - L_q) \cdot i_{sd} \cdot i_{sq}$ is negligible — the standard $i_d^* = 0$ strategy is used:

$$i_{sd}^* = 0 \qquad i_{sq}^* = T_{\text{ref}} / (p \cdot \psi_m \cdot \pi)$$

4.2 PI gains — technical-optimum tuning against the inverter delay

Per spec, gains are computed per-axis from an assumed inverter response time constant $\tau_u = 10 \mu\text{s}$ (the inverter itself is not simulated — this is purely a design parameter for gain sizing, a standard technique: choosing $K_p = L/(2\tau_u)$ places a zero that cancels the electrical pole $-R/L$, leaving a closed-loop time constant of $2\tau_u$):

$$\begin{aligned} K_{p_d} &= L_d / (2\tau_u) & K_{i_d} &= R_s / (2\tau_u) \\ K_{p_q} &= L_q / (2\tau_u) & K_{i_q} &= R_s / (2\tau_u) \end{aligned}$$

Calculated values used in the notebooks, with $R_s = 18 \text{ m}\Omega$, $L_d = 180 \mu\text{H}$, $L_q = 176 \mu\text{H}$, $\tau_u = 10 \mu\text{s}$ (so $2\tau_u = 20 \mu\text{s}$):

$$\begin{aligned} K_{p_d} &= 180\text{e-}6 / 20\text{e-}6 = 9.0 & K_{i_d} &= 18\text{e-}3 / 20\text{e-}6 = 900.0 \\ K_{p_q} &= 176\text{e-}6 / 20\text{e-}6 = 8.8 & K_{i_q} &= 18\text{e-}3 / 20\text{e-}6 = 900.0 \end{aligned}$$

4.3 Controller equations and feedforward decoupling

$$\begin{aligned} u_{sd_cmd} &= K_{p_d} \cdot (i_{sd}^* - i_{sd}) + \int K_{i_d} \cdot (i_{sd}^* - i_{sd}) dt - \omega \cdot L_q \cdot i_{sq} \\ u_{sq_cmd} &= K_{p_q} \cdot (i_{sq}^* - i_{sq}) + \int K_{i_q} \cdot (i_{sq}^* - i_{sq}) dt + \omega \cdot L_d \cdot i_{sd} + \omega \cdot \psi_m \cdot \pi \end{aligned}$$

The feedforward terms are exactly the cross-coupling / back-EMF terms from the plant model (§2.2) — added *after* the PI stage, per spec, to decouple the two current loops from the speed-dependent coupling and magnet back-EMF, so that the PI only has to reject a residual first-order error. ω and the currents used in

these expressions are the plant's real values in notebooks 01–02, and the observer's estimates in notebook 03 (Section 7).

4.4 Voltage saturation and anti-windup

$[\text{usd_cmd}, \text{usq_cmd}]$ is clamped to a circular limit $|\cdot| \leq V_{\max} = 600 \text{ V}$ before being applied (and before being handed to the observer as v_α, v_β). Each axis's integrator uses a standard back-calculation anti-windup term, $K_i \cdot \text{error} + K_p \cdot (u_{\text{sat}} - u_{\text{unsat}})$, so the integral state stops winding up while that axis is clipped. In practice this machine's voltage requirement is far below the limit even at rated conditions ($|u| \approx 108 \text{ V}$ at 70 N·m, 1500 rpm) — the saturation is a modeling safety net, not something that binds in normal operation (confirmed in notebook 03).

5. Sensorless position/speed estimation

5.1 Why back-EMF estimation

With no mechanical sensor, θ_r and ω_{mech} must be inferred from electrical quantities alone. The standard technique — used here, and described in Urbanski (2015) [see §8] — is to estimate the back-EMF vector (e_α, e_β) induced in the stationary frame, which by §2.4 points in the direction $(\cos\theta_r, \sin\theta_r)$ scaled by $\omega_r \cdot \psi_m, \text{pi}$, and recover the angle from it. The paper's own conclusion (confirmed independently here, Section 7) is that this works well away from zero speed, and has two specific, well-known weaknesses: near-zero speed (where the back-EMF vector itself vanishes) and, less commonly discussed, a **sign ambiguity at negative speed** (derived below).

5.2 Observer structure — Structure B (PI-corrected Luenberger)

The source paper describes three back-EMF estimation structures; **Structure B** (§3.B in the paper — a Luenberger observer whose plain proportional correction gains are each replaced by a full PI correction) is the one implemented here, using the machine's isotropic approximation $L_{\text{obs}} = (L_d + L_q)/2$ (valid since $L_d \approx L_q$ — the paper's $\alpha\beta$ model assumes magnetic symmetry). With $\Delta i_\alpha = i_{\alpha, \text{meas}} - \hat{i}_\alpha$ (and identically for β), and a single shared error-integral state per axis $z_\alpha = \int \Delta i_\alpha \, dt$ (since both the current- and back-EMF-channel PI corrections integrate the *same* current error, just with different gains):

$$\begin{aligned} d\hat{i}_\alpha/dt &= -(R_s/L_{\text{obs}}) \cdot \hat{i}_\alpha - (1/L_{\text{obs}}) \cdot \hat{e}_\alpha + (1/L_{\text{obs}}) \cdot v_\alpha + K_{i_P} \cdot \Delta i_\alpha + K_{i_I} \cdot z_\alpha \\ d\hat{e}_\alpha/dt &= -(K_{e_P} \cdot \Delta i_\alpha + K_{e_I} \cdot z_\alpha) \\ dz_\alpha/dt &= \Delta i_\alpha \end{aligned}$$

(β -axis: identical, mirror $\alpha \rightarrow \beta$). v_α, v_β and $i_{\alpha, \text{meas}}, i_{\beta, \text{meas}}$ are the physically measurable applied voltage and stator current, rotated into the stationary frame — in notebook 02 via the real θ_r (no mismatch yet); in notebook 03 via the controller's own estimate $\hat{\theta}$, which is where the sensorless dynamics actually bite (Section 7).

Sign note. The minus sign on the back-EMF channel ($d\hat{e}/dt = -(\dots)$, opposite in sign to the current channel) looks asymmetric but is *required* for stability — this was verified directly by Routh–Hurwitz analysis of the linearized error dynamics (not simply carried over from the source paper's typesetting, which does not make the two channels' opposite polarity obvious). Reference structure (the plain 2-state Luenberger, "Structure A") has error dynamics $d(e_i)/dt = -(R/L + K_i)e_i - (1/L)e_e$, $d(e_e)/dt = -K_e \cdot e_i$, giving characteristic polynomial $s^2 + (R/L + K_i)s + K_e/L$; both coefficients must be positive for stability, which requires the observed sign.

5.3 Observer gain derivation (pole placement)

The source paper describes the PI-corrected structure but gives no numeric tuning. Gains here are derived by placing the observer's own linearized error dynamics (a 3-state system per axis: current error $e_i = i - \hat{i}$, back-EMF error $e_e = e - \hat{e}$, and the integral state z) at a **triple real pole** $s = -\omega_n$:

$$A = \begin{bmatrix} -(R/L + K_{i_P}) & -1/L & -K_{i_I} \\ K_{e_P} & 0 & K_{e_I} \\ 1 & 0 & 0 \end{bmatrix}$$

giving characteristic polynomial:

$$s^3 + (R/L + K_{i_P}) \cdot s^2 + (K_{i_I} + K_{e_P}/L) \cdot s + K_{e_I}/L$$

Matching this to $(s + \omega_n)^3 = s^3 + 3\omega_n \cdot s^2 + 3\omega_n^2 \cdot s + \omega_n^3$ gives 2 equations for 3 remaining unknowns (the s^2 and s^0 coefficients pin down K_{i_P} and K_{e_I} uniquely; the s^1 coefficient leaves one free parameter between K_{i_I} and K_{e_P}). The remaining degree of freedom is resolved by setting K_{e_P} to the value a plain 2-state (Structure-A) observer at the same ω_n would use, which then fixes K_{i_I} :

$$\begin{aligned} K_{i_P} &= 3\omega_n - R/L & K_{e_P} &= \omega_n^2 \cdot L \\ K_{i_I} &= 2\omega_n^2 & K_{e_I} &= \omega_n^3 \cdot L \end{aligned}$$

Choosing ω_n . The observer's own model treats e_α , e_β as a slowly-varying disturbance (the pole-placement above assumes $de/dt \approx 0$ locally), but they are actually *rotating* at the electrical frequency ω_r — up to $\omega_{elec, rated} = p \cdot \omega_{mech, rated} \approx 1571$ rad/s at the 1500 rpm rated point. If ω_n is not well above that frequency, the estimated back-EMF vector picks up a **real, speed-proportional phase lag** relative to the true one. This was found empirically while building notebook 02: an initial design choice of $\omega_n = 5000$ rad/s (only $\sim 3.2\times$ the rated electrical frequency) produced a ~ 0.6 rad position error at rated speed. Sweeping ω_n confirmed the error shrinks roughly as $1/\omega_n$ without fully vanishing:

ω_n (rad/s)	position error at rated speed
5,000	~0.61 rad
20,000	~0.16 rad
40,000	~0.08 rad
60,000	~0.056 rad
80,000 (final)	~0.043 rad
100,000	~0.035 rad

$\omega_n = 80,000$ rad/s (~51× the rated electrical frequency) was chosen as the practical balance — still well-conditioned for `scipy.integrate.solve_ivp`, and comfortably resolves the fast dynamics without an excessive step-size penalty.

Calculated values used in the notebooks, with $\omega_n = 80,000$ rad/s, $R = R_s = 18$ mΩ, $L = L_{\text{obs}} = (L_d + L_q)/2 = 178$ μH (so $R/L \approx 101.12$ rad/s):

```

Ki_P = 3×80,000 - 101.12      ≈ 239,899
Ke_P = 80,000² × 178e-6       = 1,139,200
Ki_I = 2×80,000²              = 12,800,000,000 (1.28×10¹⁰)
Ke_I = 80,000³ × 178e-6       ≈ 91,136,000,000 (≈9.114×10¹⁰)

```

The residual ~0.04 rad bias that remains even here is **not further tuning error** — it is traced to the observer's isotropic L_{obs} approximation of the real, slightly salient ($L_d \neq L_q$) machine, which acts as a disturbance rotating at the electrical frequency; pure P+I correction reduces this (as gain/bandwidth increases) but cannot fully cancel a disturbance at a specific nonzero frequency — a standard sinusoidal-disturbance-rejection limit, not a bug.

5.4 Position/speed extraction — type-2 PLL

The raw angle is recovered from the estimated back-EMF vector's direction:

```
θraw = atan2(-êα, êβ)
```

(the paper's own alternative — computing speed directly from the back-EMF magnitude — is explicitly shown in the source paper to be unreliable at low speed, and our schedule crosses zero speed, so it is not used here.) θ_{raw} is fed into a standard type-2 PLL for a smooth $\hat{\theta}$, $\hat{\omega}$:

```

Kp_pll = 2·ζ·ωn,pll      Ki_pll = ωn,pll²

θ̇pll = ωpll + Kp_pll·sin(θraw - θpll)
ω̇pll = Ki_pll·sin(θraw - θpll)

```

using $\sin(\cdot)$ of the angle difference (rather than a plain subtraction) as a standard, wrap-safe phase-error signal. $\omega_{n,p11} = 15,000 \text{ rad/s}$ (also retuned upward from an initial 1,000 rad/s for the same electrical-frequency-margin reason as the observer itself — the PLL is type-2, so it has zero steady-state error for a pure ramp/constant-speed input in principle, but still needs settling-time margin) and $\zeta = 0.707$. $\hat{\theta} = \theta_{p11}$, $\hat{\omega} = \omega_{p11}$ (electrical); $\omega_{mech} = \omega_{p11}/p$.

Calculated values used in the notebooks, with $\omega_{n,p11} = 15,000 \text{ rad/s}$, $\zeta = 0.707$:

$$\begin{aligned} Kp_{p11} &= 2 \times 0.707 \times 15,000 = 21,210 \\ Ki_{p11} &= 15,000^2 = 225,000,000 \quad (2.25 \times 10^8) \end{aligned}$$

6. Model and gain parameters (reference table)

6.1 Machine and load parameters

Symbol	Value	Source
R_s	18 mΩ	machine datasheet
L_d	180 μH	machine datasheet
L_q	176 μH	machine datasheet
ψ_m (raw)	0.053 Wb	machine datasheet
$\psi_{m,pi} = \sqrt{3/2} \cdot \psi_m$	≈0.064911 Wb	derived (§2.1)
p (pole pairs)	10	machine datasheet
J	0.05 kg·m ²	user-selected (not in literature)
B	0	assumption — only the fan load damps rotation
Rated point	11 kW @ 1500 rpm	user spec
$\omega_{mech,rated}$	≈157.080 rad/s	derived
$\omega_{elec,rated} = p \cdot \omega_{mech,rated}$	≈1570.796 rad/s	derived
T_{rated}	≈70.028 N·m	derived
k_{fan}	≈2.838×10 ⁻³ N·m/(rad/s) ²	derived (§3)

6.2 Current-loop design

Symbol	Value
<code>tau</code>	10 μ s
<code>Kp_d</code>	9.0
<code>Ki_d</code>	900.0
<code>Kp_q</code>	8.8
<code>Ki_q</code>	900.0
<code>Vmax</code>	600 V

6.3 Observer (Structure B) and PLL

Symbol	Value
<code>L_obs = (Ld+Lq)/2</code>	178 μ H
<code>wn (observer)</code>	80,000 rad/s
<code>Ki_P</code>	$\approx$ 239,899
<code>Ki_I</code>	1.28×10^{10}
<code>Ke_P</code>	1,139,200
<code>Ke_I</code>	$\approx 9.114 \times 10^{10}$
<code>wn_pll</code>	15,000 rad/s
<code>zeta (PLL damping)</code>	0.707
<code>Kp_pll</code>	21,210
<code>Ki_pll</code>	2.25×10^8

All of the above are computed by formula in `pmsm_common.py` (`current_pi_gains` , `observer_gains` , `pll_gains`) from the handful of top-level constants — not hardcoded twice — so the derivations above and the executed notebooks are guaranteed consistent.

7. Known limitations and observed behavior

Two distinct weaknesses of pure back-EMF-angle estimation are visible in this build — both expected from the literature, documented rather than engineered around:

7.1 Near-zero speed

The back-EMF vector magnitude is `|e| = wr * psi_m_pi` , which is small or zero near standstill — so `theta_raw = atan2(-e_alpha, e_beta)` carries little to no information there. This is the low-speed weakness the source paper is

centrally about (its own conclusion: back-EMF methods are unreliable "for speed range below single revolution per second"). It affects the very start of the run (drive begins at rest) and the moment the drive passes through zero speed during the $-20 \text{ N}\cdot\text{m}$ reversal.

7.2 Sign ambiguity at negative speed (found while building this project, not initially expected)

This one is exact and deterministic, not just a low-magnitude noise problem, and follows directly from §2.4: $(\hat{e}_\alpha, \hat{e}_\beta) = \omega_r \cdot \psi_m \cdot \pi \cdot (\cos\theta_r, \sin\theta_r)$. For $\omega_r < 0$, this is a *negative* scalar times $(\cos\theta_r, \sin\theta_r)$ — indistinguishable, from the vector alone, from a *positive* magnitude at angle $\theta_r + \pi$. So:

$$\theta_{\text{raw}} = \text{atan2}(-\hat{e}_\alpha, \hat{e}_\beta) = \theta_r + \pi \quad \text{whenever } \omega_r < 0$$

deterministically, for as long as the speed stays negative — not a transient loss of lock. This was confirmed exactly in notebook 02 (passive observer: position error sits at a flat π throughout the $-20 \text{ N}\cdot\text{m}$ level, recovering the instant speed goes positive again) and has real consequences once notebook 03 lets the controller act on it:

- **During the reversal:** the controller's estimated currents come out negated relative to the real ones ($\hat{\text{Park}}(\theta)$ at $\theta \approx \theta_r + \pi$ flips both axes), so it chases the wrong sign of error — producing a violent, sustained oscillation (T_e swings to roughly $\pm 220 \text{ N}\cdot\text{m}$, over $3\times$ rated torque) rather than a quiet mistracking, and driving the adaptive ODE solver into a much smaller step size (chaotic dynamics, not merely a slow transient).
- **After the reversal:** the drive does **not** recover on its own. $\hat{\theta} = \theta_r + \pi$ turns out to be a genuine, self-consistent equilibrium of the closed loop: with the frame permanently flipped, the controller's estimated $\hat{i}_q$ is $-i_{q,\text{real}}$, so driving that toward $+i_q^*$ drives the real i_q toward $-i_q^*$ — a commanded $+50 \text{ N}\cdot\text{m}$ settles into a stable response that looks like $-50 \text{ N}\cdot\text{m}$ was commanded instead (confirmed in notebook 03: $T_e \approx -50 \text{ N}\cdot\text{m}$ at the end of the run, $\hat{\omega}$ tracking ω_{mech} accurately in *rate* the whole time, just phase-shifted by exactly π). Nothing later in the schedule dislodges it.

Real sensorless drives resolve this with an independent sign-of-speed reference (e.g. from the commanded i_q , or a remembered pre-reversal state) or a re-synchronization/realignment procedure — genuinely useful future work, but out of scope for this project, which is specifically about the plain back-EMF-observer structure from the source literature.

8. References

- "RL_Modelo PMSM.pdf" (internal reference document) — dq machine model derivation (generator, power-invariant transform), re-derived to motor convention in Section 2.
- "RL_Parametros_PMSM.png" (internal reference document) — machine parameters (Section 6.1).
- Urbanski, K., "Estimation of Back EMF for PMSM at Low Speed Range", MM Science Journal, March 2015 — $\alpha\beta$ back-EMF observer structures (Structure A/B), Section 5.

- `pmsm_common.py` (in this repository) — the single source of truth for every constant and gain formula quoted above.